

# Hands-on Session – Part II



# Compare *dist* on CPU and GPU

- This exercise still uses the Google Colab environment introduced in the 1<sup>st</sup> hands-on session
  - Using *cuda\_Colab\_part2.ipynb*
  - *Separate codes can be found in* [dist\\_gpu.cu](#), [dist\\_cpu.c](#)
- First compile a CPU version of the 1D dist program introduced in Lecture 2
- Then compile the GPU version
- Run two batches of tests using CPU and GPU versions
- Plot their performance difference



# 1 compile the CPU version of the 1D dist program



cuda\_Colab\_part2.ipynb ☆

File Edit View Insert Runtime Tools Help [All changes saved](#)

+ Code + Text

## Write the CPU version of the 1D *dist* program in Lecture 2

```
%%writefile dist_cpu.c
#include <sys/time.h>
#include <stdio.h>
#include <stdlib.h>
#include <math.h> //Include standard math library containing sqrt.

// Timer
double get_cpu_millisecond() {
    struct timeval tp;
    gettimeofday(&tp, NULL);
    return ((double)tp.tv_sec*1.e6 + (double)tp.tv_usec);
}

// A scaling function to convert integers 0,1,...,N-1 to evenly spaced floats
```




# 1 compile the CPU version of the 1D dist program

A screenshot of a Jupyter Notebook interface. The top bar shows the Colab logo, the file name 'cuda\_Colab\_part2.ipynb', and a star icon. Below this is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Runtime', 'Tools', 'Help', and a link 'All changes saved'. The left sidebar contains icons for a menu, search, and a variable '{x}'. The main area has a tab '+ Code' and a section header 'Compile the CPU version code'. Below the header is a code cell with the following text:

```
[ ] !gcc -o dist_cpu dist_cpu.c -lm
    !ls
```



# 2 compile the GPU version of the 1D dist program

 cuda\_Colab\_part2.ipynb 

File Edit View Insert Runtime Tools Help

+ Code + Text

▼ Write the GPU version of the 1D *dist* program in Lecture 2



```
%%writefile dist_gpu.cu
#include <sys/time.h>
#include <stdio.h>
#include <math.h> //Include standard math library containing sqrt.

#define TPB 128    // Specify threads per block

// Timer
double get_cpu_millisecond() {
    struct timeval tp;
```



# 2 compile the GPU version of the 1D dist program

The screenshot shows a JupyterLab interface with a file browser on the left and a code editor on the right. The file browser shows a directory structure with files like `dist_cpu.c`, `dist_gpu.cu`, `sample_data`, `vectorAdd`, `vectorAdd_naive.cu`, and `vectorAdd_v2.cu`. The code editor has a tab for `cuda_Colab_part2.ipynb` and contains two code cells. The first cell contains the commands `!nvcc -o dist_gpu dist_gpu.cu` and `!ls`. The second cell contains the command `!./dist_gpu 1024`. The output of the second cell is `36.00,`.

```
cuda_Colab_part2.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

▼ Compile the GPU version

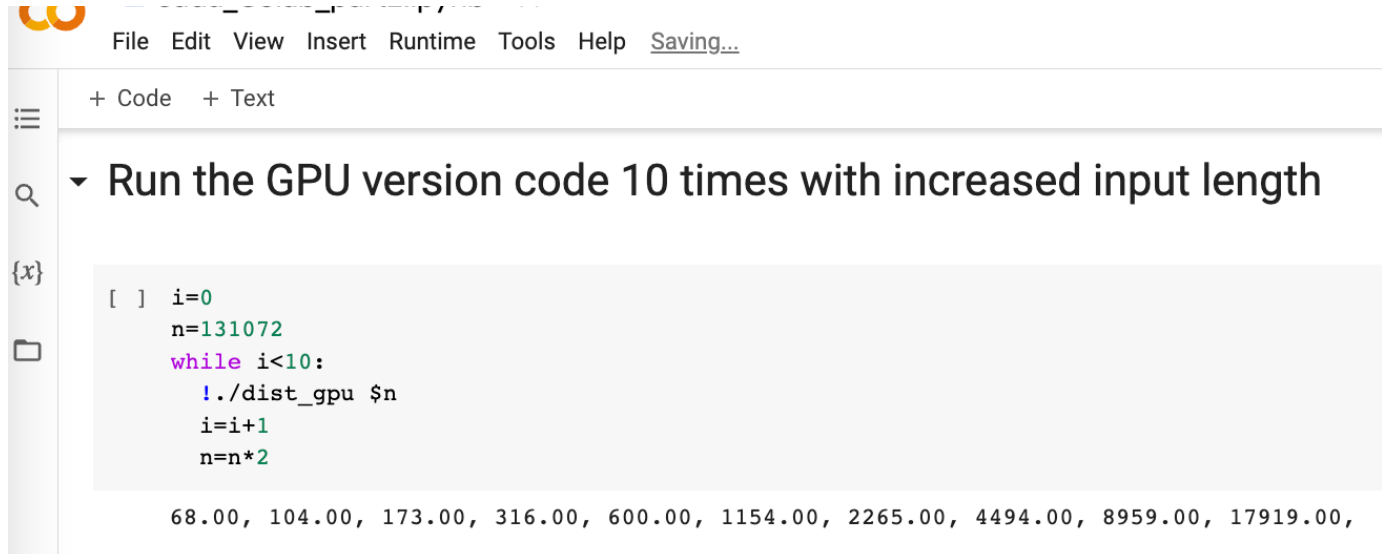
!nvcc -o dist_gpu dist_gpu.cu
!ls

dist_cpu    dist_gpu    sample_data  vectorAdd_naive  vectorAdd_v2.cu
dist_cpu.c  dist_gpu.cu  vectorAdd    vectorAdd_naive.cu

!./dist_gpu 1024

36.00,
```

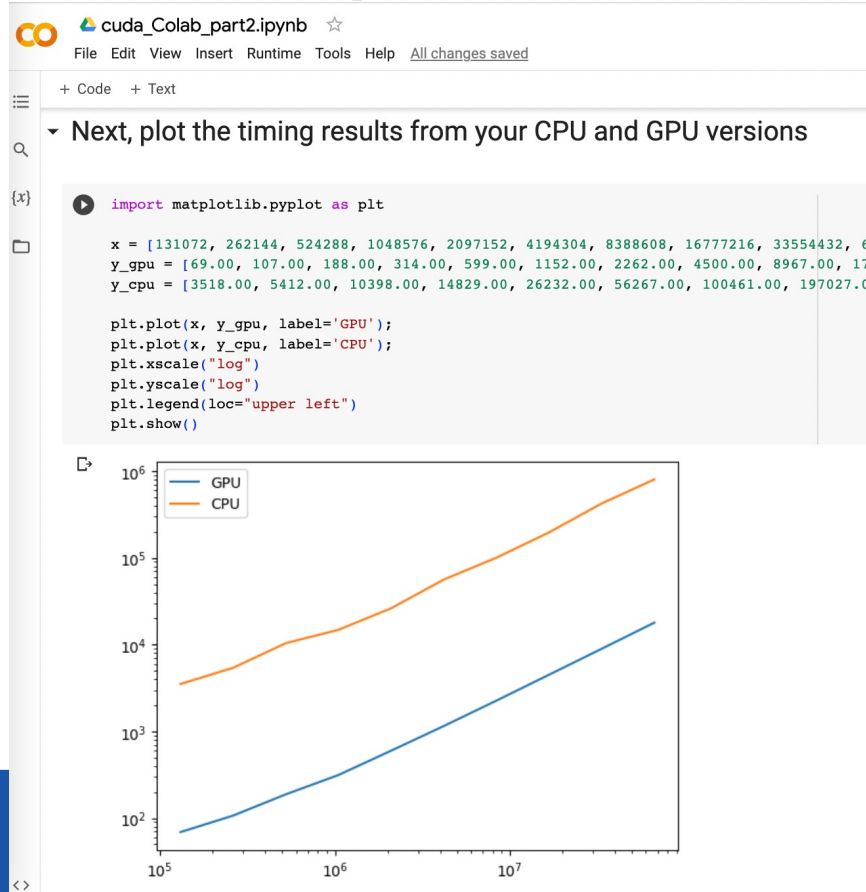
# 3 run a batch of tests using different input

A screenshot of a code editor interface. The top menu bar includes "File", "Edit", "View", "Insert", "Runtime", "Tools", "Help", and "Saving...". Below the menu bar, there are tabs for "+ Code" and "+ Text". The main editor area displays a code snippet with a title "Run the GPU version code 10 times with increased input length". The code is as follows:

```
[ ] i=0
n=131072
while i<10:
    !./dist_gpu $n
    i=i+1
    n=n*2
```

Below the code, the output of the program is shown: "68.00, 104.00, 173.00, 316.00, 600.00, 1154.00, 2265.00, 4494.00, 8959.00, 17919.00,". The editor has a sidebar on the left with icons for a menu, search, variables, and a file explorer.

# Plot and Compare GPU and CPU performance







# Optimize a naïve 2D matrix-multiplication

- This exercise still uses the Google Colab environment introduced in the 1<sup>st</sup> hands-on session
  - Using `cuda_Colab_part2.ipynb`
  - *Separate codes can be found in [sgemm\\_naive.cu](#)*
- First compile and run the naïve version on GPU
- Run two batches of tests using CPU and GPU versions
- Plot their performance difference



# 1 Compile and run the naïve version on GPU



cuda\_Colab\_part2.ipynb ☆

File Edit View Insert Runtime Tools Help

+ Code + Text



## Write the GPU version of a naive 2D matrix multiplication



```
%%writefile sgemm_naive.cu
```

```
#include <stdio.h>
```

```
#include <sys/time.h>
```

```
#define DataType double
```

```
// Compute C = A * B
```

```
// Sgemm stands for single precision general matrix-matrix multiply
```

```
global void gemm(DataType *A, DataType *B, DataType *C, int numARows,
```



## 2 Run the code with different input

The screenshot shows a JupyterLab window titled 'cuda\_Colab\_ex1.ipynb'. The interface includes a top menu bar with 'File', 'Edit', 'View', 'Insert', 'Runtime', 'Tools', and 'Help', along with a status bar indicating 'All changes saved'. On the left, a sidebar contains icons for a file explorer, search, and a variable inspector. The main area displays a code cell with the following content:

```
[ ] !./sgemm_naive 128 128 128
```

Below the code cell, the output is shown:

```
Input matrix dim (128 x 128) (128 x 128) (128 x 128)
valid
```

# 3 Profile the code with Nvprof

▸ next, profiling with nvprof



```
!nvprof ./sgemm_naive 128 512 512
```



```
Input matrix dim (128 x 512) (512 x 512) (128 x 512)
```

```
==12051== NVPROF is profiling process 12051, command: ./sgemm_naive 128 512 512
```

```
valid
```

```
==12051== Profiling application: ./sgemm_naive 128 512 512
```

```
==12051== Profiling result:
```

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	65.22%	823.02us	1	823.02us	823.02us	823.02us	823.02us	gemm(double*, double*, do
	30.86%	389.43us	2	194.72us	45.695us	343.74us	343.74us	[CUDA memcpy HtoD]
	3.92%	49.407us	1	49.407us	49.407us	49.407us	49.407us	[CUDA memcpy DtoH]
API calls:	98.82%	205.10ms	3	68.366ms	3.5130us	205.02ms	205.02ms	cudaMalloc
	0.98%	2.0247ms	3	674.91us	132.41us	1.2980ms	1.2980ms	cudaMemcpy
	0.12%	239.84us	3	79.947us	15.321us	118.01us	118.01us	cudaFree
	0.05%	112.57us	101	1.1140us	143ns	47.178us	47.178us	cuDeviceGetAttribute
	0.01%	28.986us	1	28.986us	28.986us	28.986us	28.986us	cudaLaunchKernel
	0.01%	25.681us	1	25.681us	25.681us	25.681us	25.681us	cuDeviceGetName
	0.01%	11.400us	1	11.400us	11.400us	11.400us	11.400us	cuDeviceGetPCIBusId
	0.00%	2.9870us	2	1.4930us	1.1030us	1.8840us	1.8840us	cuDeviceGet
	0.00%	2.7060us	3	902ns	228ns	2.1710us	2.1710us	cuDeviceGetCount
	0.00%	589ns	1	589ns	589ns	589ns	589ns	cuDeviceTotalMem
	0.00%	493ns	1	493ns	493ns	493ns	493ns	cuModuleGetLoadingMode
	0.00%	236ns	1	236ns	236ns	236ns	236ns	cuDeviceGetUuid



# 2D Matrix Multiplication – Exercise

For a matrix A of (128x128) and B of (128x128):

- Explain how many CUDA threads and thread blocks are used in your tests.
- Profile your program with nvprof. What are the top 3 activities in time?



## 2D Matrix Multiplication – Exercise

For a matrix A of (511x1023) and B of (1023x4094):

- Did your program still work? If not, what changes did you make?
- Explain how many CUDA threads and thread blocks you used.

# 2D Matrix Multiplication – Exercise

Optimize the naïve implementation of with shared memory and tiles

```
__global__ void gemm(DataType *A, DataType *B, DataType *C, int numRows,
                    int numAColumns, int numBRows, int numBColumns){
    /// Insert code to implement matrix multiplication here
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    if (row < numRows && col < numBColumns) {
        DataType sum = 0;
        for (int ii = 0; ii < numAColumns; ii++) {
            sum += A[row * numAColumns + ii] * B[ii * numBColumns + col];
        }
        C[row * numBColumns + col] = sum;
    }
}
```

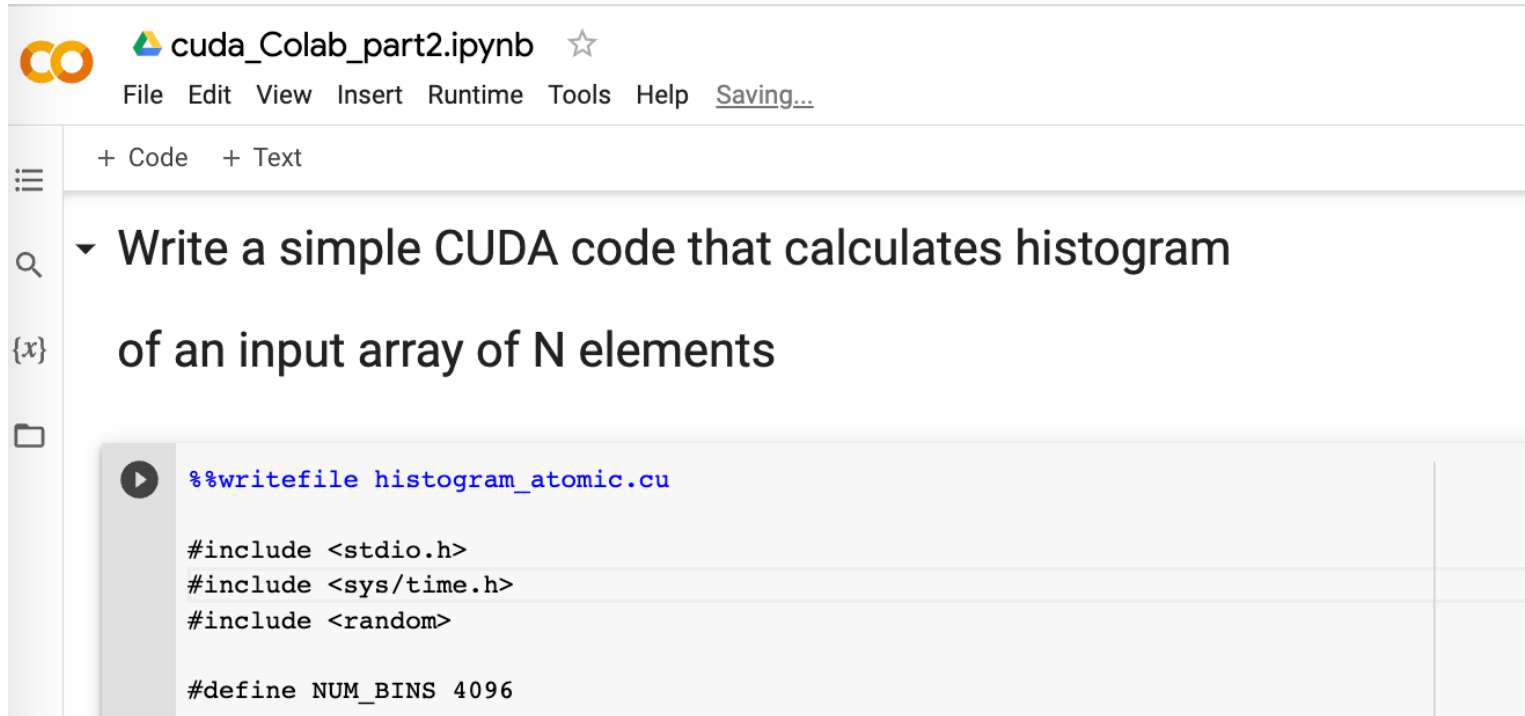


# Calculating histogram using atomic operation

- This exercise still uses the Google Colab environment introduced in the 1<sup>st</sup> hands-on session
  - Using `cuda_Colab_part2.ipynb`
  - *Separate codes can be found in* [Histogram\\_atomic.cu](#)
- First compile and run the naïve GPU version
- Run a batch of tests using increased array
- Plot the output histogram



# 1 Compile the naïve histogram code



cuda\_Colab\_part2.ipynb ☆

File Edit View Insert Runtime Tools Help Saving...

+ Code + Text

Write a simple CUDA code that calculates histogram of an input array of N elements

```
%%writefile histogram_atomic.cu

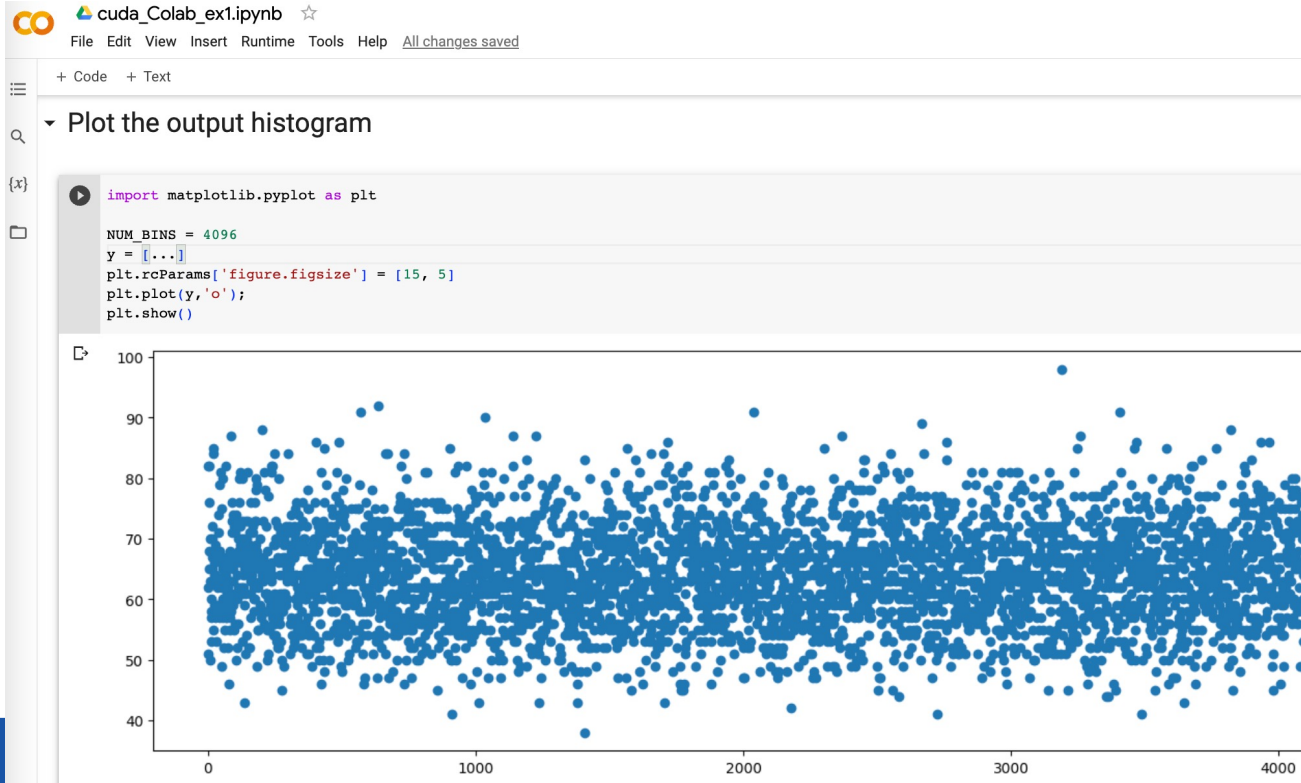
#include <stdio.h>
#include <sys/time.h>
#include <random>

#define NUM_BINS 4096
```



18

# Plot and Compare the output histogram





# Histogram – Exercise

For an array of length 1048576:

- How many global reads are performed?
- How many atomic operations are used?



# Histogram – Exercise

Optimize the naïve implementation to reduce the number of atomic operations used

```
__global__ void histogram_kernel(unsigned int *input, unsigned int *bins,
                                unsigned int num_elements,
                                unsigned int num_bins) {

    unsigned int tid = blockIdx.x * blockDim.x + threadIdx.x;

    // Privatized bins
    extern __shared__ unsigned int bins_s[];
    for (unsigned int binIdx = threadIdx.x; binIdx < num_bins;
         binIdx += blockDim.x) {
        bins_s[binIdx] = 0;
    }
    __syncthreads();

    // Histogram
    for (unsigned int i = tid; i < num_elements; i += blockDim.x * gridDim.x) {
        atomicAdd(&(bins_s[input[i]]), 1);
    }
    __syncthreads();

    // Commit to global memory
    for (unsigned int binIdx = threadIdx.x; binIdx < num_bins;
         binIdx += blockDim.x) {
        atomicAdd(&(bins[binIdx]), bins_s[binIdx]);
    }
}
```

# Q & A