

Hands-on: part IV

Connect to Dardel - Kerberos

Before connecting to Dardel, the local environment of your laptop must be configured.

Detailed instructions can be found on <https://www.pdc.kth.se/support/documents/login/configuration.html>.

- If you are using **Mac**, you need to follow a special set of instructions here: https://www.pdc.kth.se/support/documents/login/mac_login.html.

- If you are using **Windows**, we recommend using the WSL2 to use a fully-fledged Ubuntu: https://www.pdc.kth.se/support/document/login/windows_login.html#wsl-approach.

To configure Kerberos on your local machine, you need to do the following steps:

- 1) Create .ssh folder if it does not already exist in your home folder.
- 2) Create a file called krb5.conf in .ssh with the following content





Connect to Dardel - SSH

1. See <https://www.pdc.kth.se/support/documents/login/configuration.html>
2. Create a file in `.ssh` called `config` with the following content

```
# Hosts we want to authenticate to with Kerberos
Host *.kth.se *.kth.se.
# User authentication based on GSSAPI is allowed
GSSAPIAuthentication yes
# Key exchange based on GSSAPI may be used for server authentication
GSSAPIKeyExchange yes
# Hosts to which we want to delegate credentials. Try to limit this to
# hosts you trust, and were you really have use for forwarded tickets.
Host *.csc.kth.se *.csc.kth.se. *.nada.kth.se *.nada.kth.se. *.pdc.kth.se *.pdc.kth.se.
# Forward (delegate) credentials (tickets) to the server.
GSSAPIDelegateCredentials yes
# Prefer GSSAPI key exchange
PreferredAuthentications gssapi-keyex,gssapi-with-mic
# All other hosts
Host *
```

2. Set the correct permission on the file

```
$ chmod 644 ~/.ssh/config
```

3. Connect

```
$ ssh YourUsername@dardel.pdc.kth.se
```



Use Slurm to Launch Your Simulation on Computing Nodes

- Slurm allows us to launch our simulation on the computing nodes
- Two ways:
 - **Interactive mode:** ask for a resource, and as soon it becomes available, we can launch your simulations interactively

```
$ salloc --nodes=1 -t 00:30:00 -A edu23.summer -p gpu --ntasks-per-node=1 --cpus-per-task=1
```

```
$ srun -n 1 ./hello.out
```

- **Batch mode:** put your simulation into a queue



Login Node



Comp.
Node

Comp.
Node

Login Node



Comp.
Node



Comp.
Node

Profiling with rocprof



Profiling Reduction with rocprof

Run rocprof with `--stats` for a summary in `results.stats.csv`

```
4_reduction> srunk -n 1 rocprof --stats ./reduction
RPL: on '230811_145940' from '/cfs/klemming/root/rocm/opt/rocm-5.3.3' in './amd_part2/4_reduction'
RPL: profiling './reduction'
RPL: input file ''
RPL: output dir '/tmp/rpl_data_230811_145940_117170'
RPL: result dir '/tmp/rpl_data_230811_145940_117170/input_results_230811_145940'
Usage: ./reduction num_of_elems
using default value: 52428800
ARRAYSIZE: 52428800
Array size: 200 MB
ROCProfiler: input from "/tmp/rpl_data_230811_145940_117170/input.xml"
  0 metrics
The average performance of reduction is 396.195 GBytes/sec
VERIFICATION: result is CORRECT

ROCProfiler: 20 contexts collected, output directory /tmp/rpl_data_230811_145940_117170/input_results_230811_145940
File '/amd_part2/4_reduction/results.csv' is generating
File '/amd_part2/4_reduction/results.stats.csv' is generating
```



Profiling Reduction with rocprof

Run rocprof with a selected list of hardware counters specified in my_counters.txt

```
4_reduction> cat my_counters.txt
pmc: Wavefronts VALUInsts VFetchInsts VWriteInsts VALUUtilization VALUBusy WriteSize

4_reduction> srun -n 1 rocprof -i my_counter.txt ./reduction
...

ROCProfiler: input from "/tmp/rpl_data_230810_164433_16749/input0.xml"
gpu_index =
kernel =
range =
7 metrics
  Wavefronts, VALUInsts, VFetchInsts, VWriteInsts, VALUUtilization, VALUBusy, WriteSize

The average performance of reduction is 91.2052 GBytes/sec
VERIFICATION: result is CORRECT

ROCProfiler: 20 contexts collected, output directory /tmp/rpl_data_230810_164433_16749/input0_results_230810_164433
File '4_reduction/my_counter.csv' is generating
```



Profiling Reduction with rocprof

Inspect the output *my_counters.csv*

```
4_reduction> cat my_counters.csv  
  
Index,KernelName,gpu-id,queue-id,queue-  
index,pid,tid,grd,wgr,lds,scr,arch_vgpr,accum_vgpr,sgpr,wave_size,sig,obj,Wavefronts,VALUInsts,VFetchInsts,VWriteInsts,VALUUtiliza  
tion,VALUBusy,WriteSize  
.....  
..
```

Question: post your output to our Slack channel

Matrix Multiplication in HIP



Matrix Multiplication in HIP

1. Compile and run the example matrix multiplication code on the Dardel supercomputer

```
5_MatrixMultiplication> make  
hipcc --offload-arch=gfx90a -I./include -c -o MatrixMultiplication.o MatrixMultiplication.cpp  
hipcc MatrixMultiplication.o -o MatrixMultiplication
```



Matrix Multiplication in HIP

2. Compile and run the example matrix multiplication code on the Dardel supercomputer

```
5_MatrixMultiplication> srun -n 1 ./MatrixMultiplication -h
Usage
-h, --help                Display this information
-q, --quiet               Quiet mode. Suppress all text output.
-e, --verify              Verify results against reference implementation.
-t, --timing               Print timing.
-v, --version             AMD APP SDK version string.
-d, --deviceId            Select deviceId to be used[0 to N-1 where N is number devices available].
-x, --height0             height of matrix A
-y, --width0              width of matrix A and Height of matrix B
-z, --width1              width of matrix B
-b, --blockSize           Use local memory of dimensions blockSize x blockSize
-i, --iterations          Number of iterations for kernel execution
    --eAppGflops           Prints GFLOPS calculated from transfer + kernel time
```



Matrix Multiplication in HIP

2. Run the program with different input size for A and B

```
5_MatrixMultiplication> srun -n 1 ./MatrixMultiplication -q -t -x 1024 -y 1024 -z 1024 -i 10 --timing --eAppGflops
```

Question 1: what is the maximum GFlops achieved?

Use rocBLAS library on AMD GPU

Use rocBLAS library on AMD GPU

1. compile and run the BLAS level 3 matrix-matrix multiplication example. Remember to load rocm module on the Dardel supercomputer

```
6_rocBLAS_gemm> module load rocm/5.3.3

6_rocBLAS_gemm> make
hipcc -std=c++14 -I./common -isystem/cfs/klemming/root/rocm/opt/rocm-5.3.3/include -
isystem/cfs/klemming/root/rocm/opt/rocm-5.3.3/include -g -Ofast -march=native -Wall -c
gemm.cpp -o gemm.o
hipcc -std=c++14 -I./common -isystem/cfs/klemming/root/rocm/opt/rocm-5.3.3/include -
isystem/cfs/klemming/root/rocm/opt/rocm-5.3.3/include -g -Ofast -march=native -Wall -c
common/ArgParser.cpp -o common/ArgParser.o
hipcc gemm.o ./common/ArgParser.o -L/cfs/klemming/root/rocm/opt/rocm-5.3.3/lib -
L/cfs/klemming/root/rocm/opt/rocm-5.3.3/lib -lrocblas -Wl,-
rpath=/cfs/klemming/root/rocm/opt/rocm-5.3.3/lib -Wl,-rpath=/cfs/klemming/root/rocm/opt/rocm-
5.3.3/lib -lm -lpthread -lstdc++ -o 6_rocBLAS_gemm
```

Use rocBLAS library on AMD GPU

1. Try with different input and check the output (you might want to uncomment the CPU part to check the correctness of the results)
- 2.

```
6_rocBLAS_gemm
--K <value>           Matrix/vector dimension
--M <value>           Matrix/vector dimension
--N <value>           Matrix/vector dimension
--alpha <value>       Alpha scalar
--beta <value>        Beta scalar

_rocBLAS_gemm> srun -n 1 ./6_rocBLAS_gemm --K 128 --M 128 --N 128 --alpha 0.25 --beta 0.75
```



Use rocBLAS library on AMD GPU

Exercise 1: convert your code in [5_MatrixMultiplication](#) to use rocBLAS sgemm in this example.

Exercise 2: rerun the matrix multiplication using rocBLAS and compare your achieved Gflops with the original matrix multiplication implementation.

Q & A