




Introduction: GPU Architecture

Ivy Peng

Assistant Professor in Computer Science

Scalable Parallel System (ScaLab)

Department of Computer Science, KTH

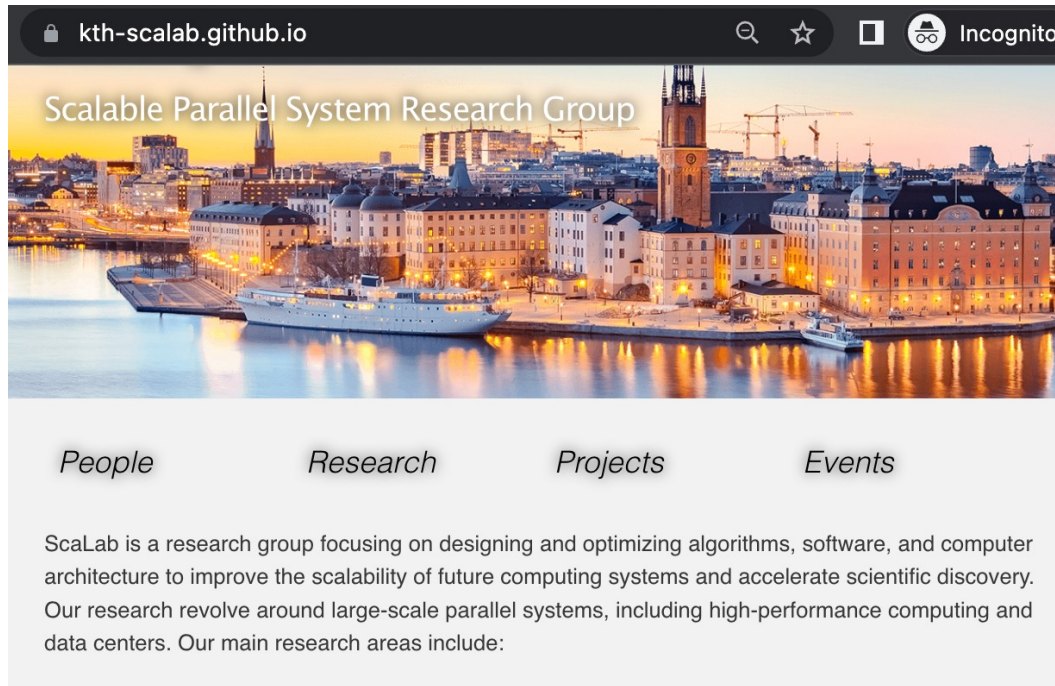




Who are we?

Scalable Parallel System (KTH-ScaLab)


- **Converged Cloud and HPC systems**
 - Kubernetes
 - Architecture
 - Workflows
- **Memory Systems**
 - Heterogeneous memory
 - Disaggregated memory
- **Heterogeneous Computing**
 - GPU
 - RISC-V
 - Quantum Computing



Intended Learning Outcomes (ILO)


At the end of this course, you will be able to:

1. Describe the **architecture** of recent AMD and Nvidia GPUs
2. Compile and run **CUDA** and **HIP** kernels on GPUs on a cluster
3. Use **profiling** tools to measure and analyze the performance of CUDA and HIP codes
4. Transform and **implement** a serial kernel into CUDA and HIP code on GPU



Course Organization: Computer Resources

- For AMD GPU: you will be given access to the PDC Dardel Supercomputer with an allocation
 - You will learn how you connect and run jobs on supercomputers as part of this course
- For Nvidia GPU: Google Colab on Google Cloud





Outline

Thursday – part 1

GPU Architecture: AMD and Nvidia
CUDA Programming
Hands-on

Thursday – part 2

CUDA Programming
Hands-on

Friday – part 3

AMD GPU
Hands-on

Friday – part 4

HIP Programming
Hands-on

Graphical Processing Unit (GPU)


GPU = specialized accelerator for processing images in video frame for display devices.

GPUs are used in game consoles, embedded systems (like systems on cars for automatic driving), computers, and supercomputers.

- Since 2012, GPUs are the main workforce for training deep-learning networks

Some important GPU vendors: **NVIDIA**, **AMD**,

...



Integrated GPU v.s. Dedicated GPU

- The main difference lies in the memory:
 - Integrated GPU shares the system memory with CPU
 - Dedicated GPU has its own memory
- Integrated GPU is often found in laptops, more power efficient, e.g., Intel HD or Iris Graphics.
- Dedicated GPUs are often removable and need more power, and provide higher performance
- In HPC, we focus on dedicated GPUs




CPUs with integrated graphics



Discrete graphics card

GPU v.s. CPU Architecture



CPU has tens of massive cores, CPU excels at irregular control-intensive work

- Lots of hardware for control, fewer ALUs

GPU has thousands of small cores, GPU excels at regular math-intensive work

- Lots of ALUs, little hardware for control


GPU Hardware Model for Nvidia GPUs

The fundamental computing entity is

- **Streaming Processor (SP) or CUDA core**

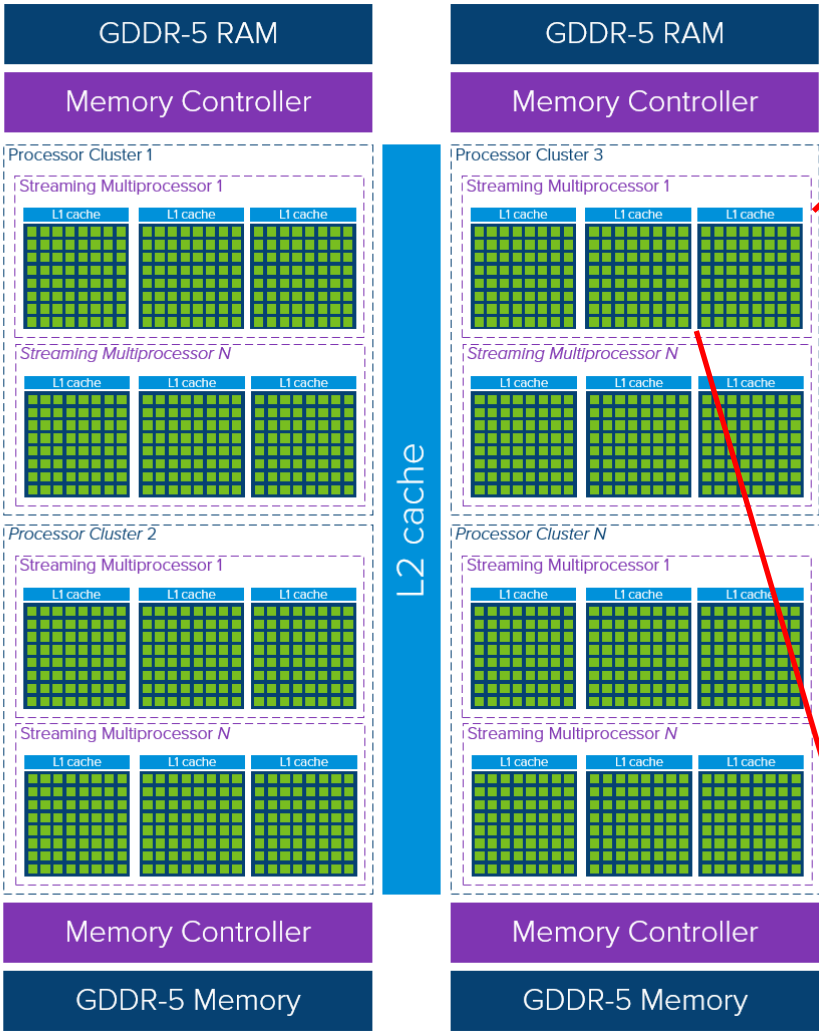
A Streaming Multiprocessor (SM):

- A collection of 8/32/192 CUDA Cores (depends on SM architecture)
- Has some fast cache shared memory
- Can synchronize

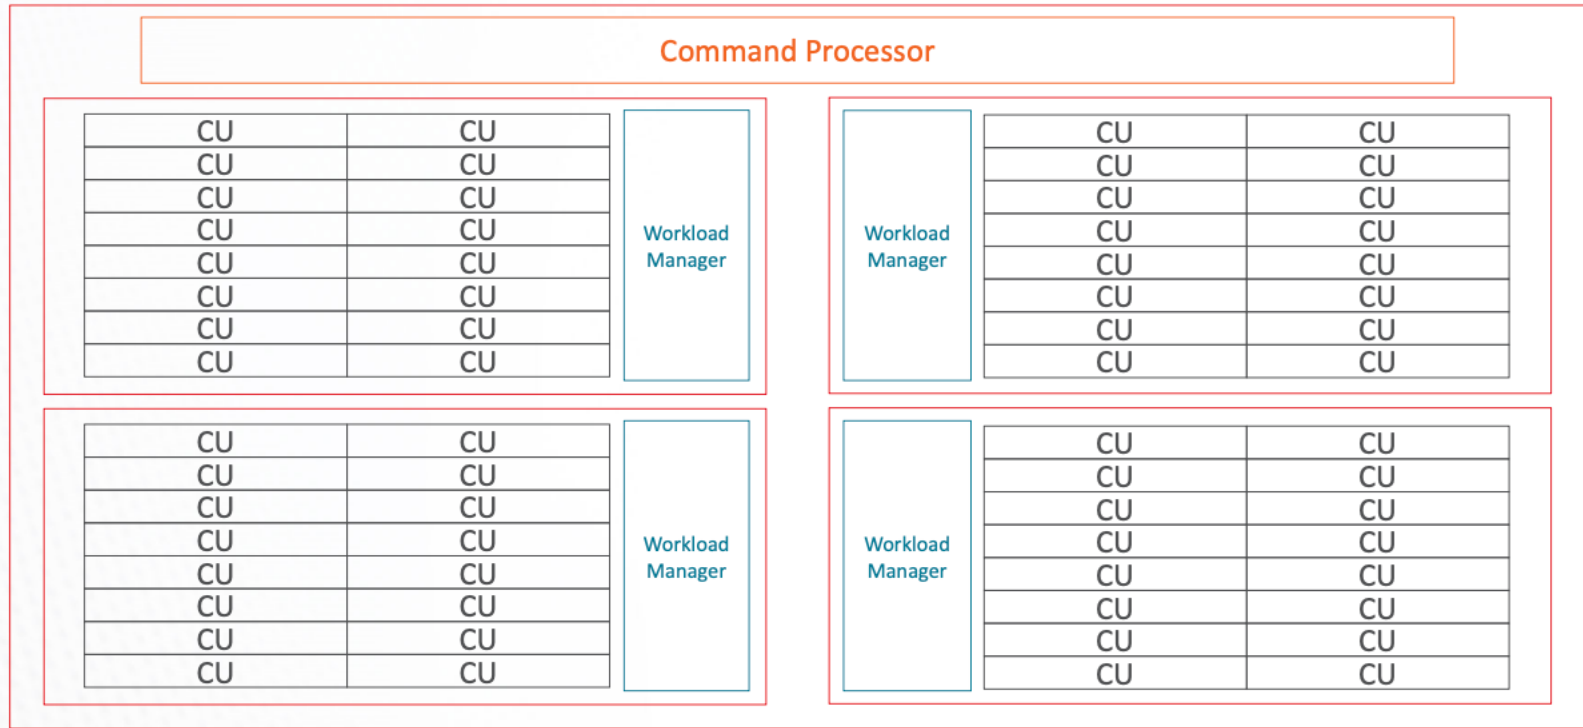


Nvidia A100 GPU

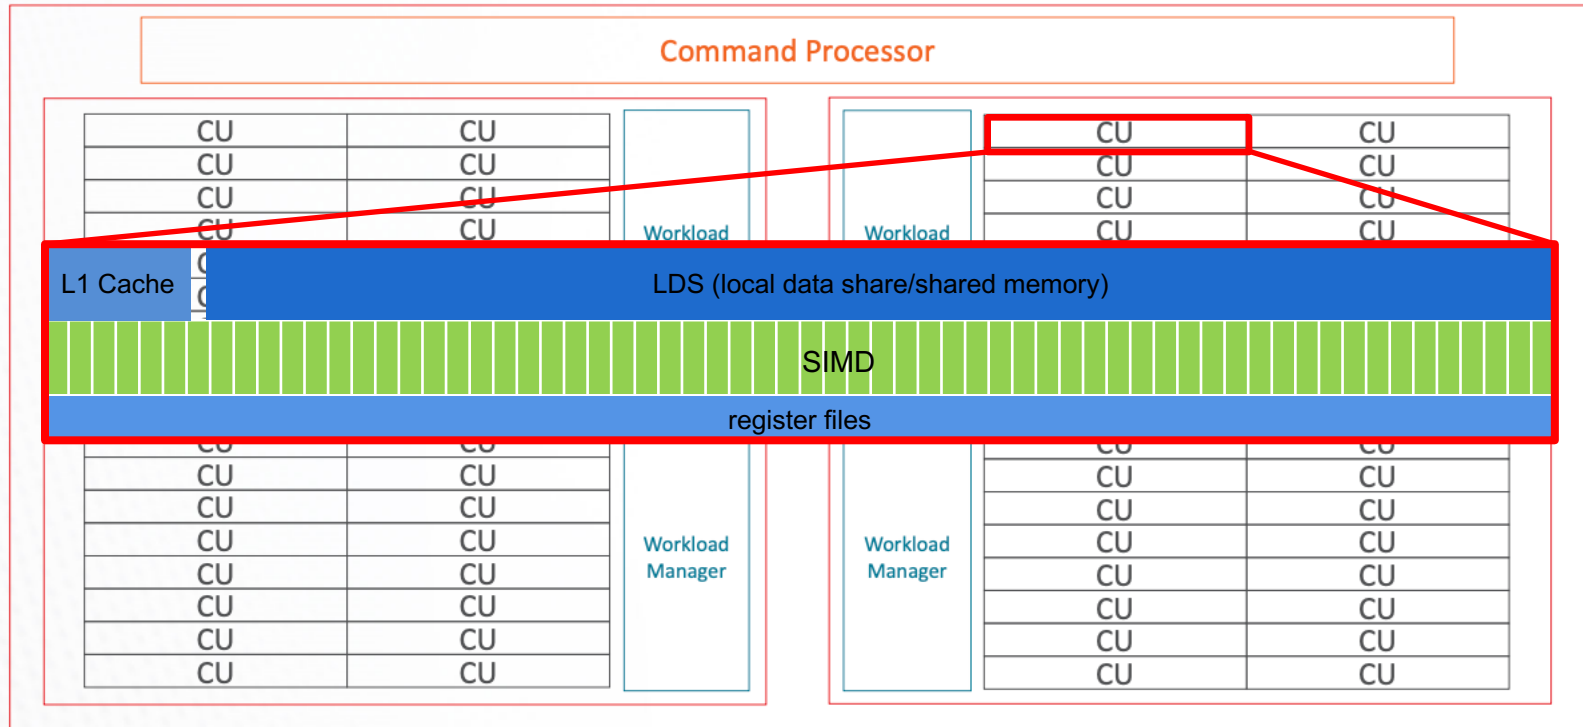
PCIe x16 3.0 host interface



AMD Graphics Core Next (GNC) Architecture



AMD Graphics Core Next (GNC) Architecture



Recent Nvidia GPU Architecture

- Nvidia Volta Architecture, tensor cores, mixed precision
 - the GA100 GPU has 128 SMs, 64 FP32 CUDA Cores/SM
- Nvidia Ampere Architecture, 3rd gen NVLink
- Nvidia Hopper Architecture

Questions: how many cores in GA100?

$$2 \times 13 \times 192 = \mathbf{4992!}$$

Questions: how many cores per node on Dardel host?


$$2 \times 64 = \mathbf{128}$$




GPU is Throughput-oriented Architecture

- GPUs focus on executing many computation in parallel to maximizing the **total throughput**
 - GPUs do not target to minimize the latency of a single task

CPU



GPU



GPU Design Motivation: Process Pixels in Parallel

1. Data parallel


- In 1080i and 1080p videos, 1920 x 1080 pixels = 2M pixels per video frame → compute intensive
- Lots of parallelism at low clock speed → power efficient

2. Computation on each pixel is independent from computation on other pixels

- No need for synchronization

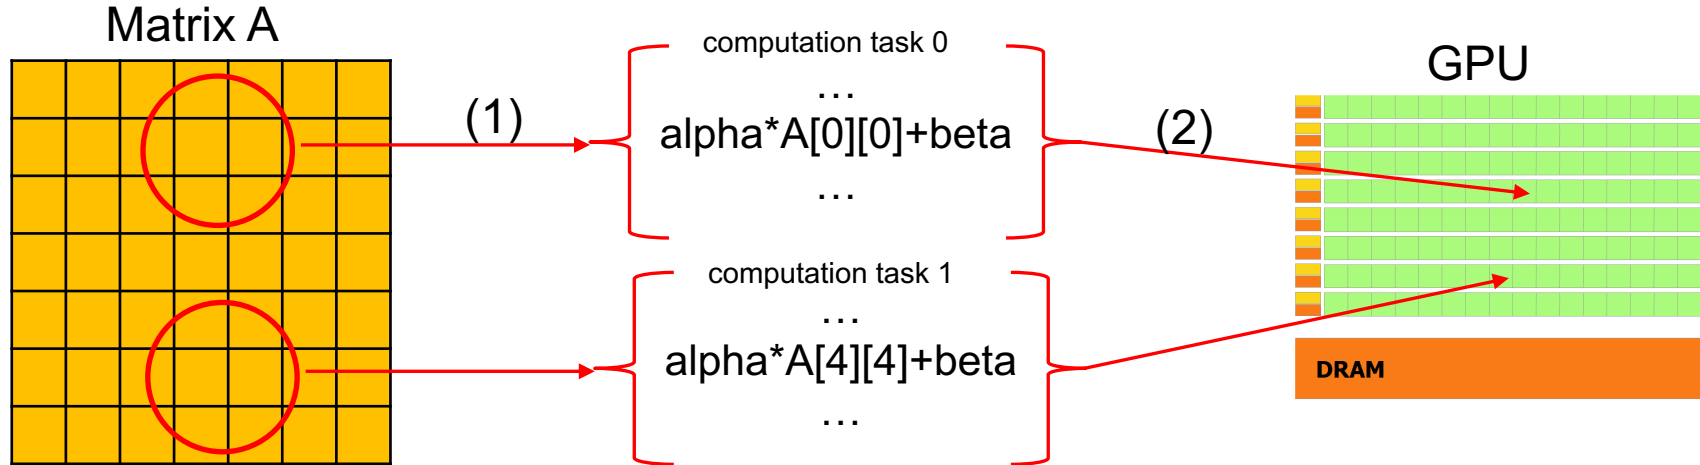
3. Good data-locality = access to data is regular

- No need for large caches




Mapping computation tasks to GPU hardware

- (1) Define computation tasks – done by you
- (2) Schedule computational tasks on GPU cores – done by CUDA runtime
 - Intuitively, higher throughput when more GPU cores are busy – **How?**



Hardware Multithreading

- Computation can be divided into a collection of many concurrent sequential tasks that executed across many threads
 - E.g., decompose a large matrix
- Thread can be seen as **virtualized scalar processor** with a program counter, register file and associate processor state
- Multithreading can be implemented in software (OS) or hardware
 - E.g., hyper-threading on Intel processor
- **Throughput-oriented architectures have implemented in hardware**



How to improve GPU utilization?


- Increase # of CUDA cores
 - High-end GPUs have a large number of cores
- Define computation tasks with low dependency and synchronization
 - If task 0 needs to wait for task 1, adding CUDA cores won't help
- Define many tasks to **oversubscribe**
 - Define $\#tasks \gg \#cores$



Hardware Multithreading Hides Latency

Long-latency operations of a single-thread can be hidden or covered by **ready-to-run work** from another thread, examples:

- Thread 1 cannot run because waiting data from DRAM
- Thread 2 can run because all its required operands are ready
- Switch to run thread 2 while overlapping data fetching for thread 1



Q & A